

- [10] F.M. Megahed, Y.J. Chen, B.M. Colosimo, and M.L.G. Grasso. Adapting OpenAI’s CLIP model for few-shot image inspection in manufacturing quality control. *arXiv preprint arXiv:2501.12596*, 2025.
- [11] K. Moenck, D.T. Thieu, J. Koch, and T. Schüppstuhl. Industrial language-image dataset (ILID): Adapting vision foundation models for industrial settings. *arXiv preprint arXiv:2406.09637*, 2024.
- [12] K. Song and Y. Yan. A noise robust method based on completed local binary patterns for hot-rolled steel strip surface defects. *Applied Surface Science*, 285:858–864, 2013.
- [13] B. Totino, F. Spagnolo, and S. Perri. RIAWELC: A novel dataset of radiographic images for automatic weld defects classification. *Research Square (Preprint)*, 2022.
- [14] P. Wang, S. Bai, S. Tan, S. Wang, Z. Fan, J. Bai, K. Chen, X. Liu, J. Wang, W. Ge, Y. Fan, K. Dang, M. Du, X. Ren, R. Men, D. Liu, C. Zhou, J. Zhou, and J. Lin. Qwen2-VL: Enhancing vision-language model’s perception of the world at any resolution. *arXiv preprint arXiv:2409.12191*, 2024.
- [15] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. Chi, Q.V. Le, and D. Zhou. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, volume 35, pages 24824–24837, 2022.
- [16] M.J. Yin, D. Jiang, Y. Chen, B. Wang, and C. Ling. Enhancing generalization in chain of thought reasoning for smaller models. *arXiv preprint arXiv:2501.09804*, 2025.
- [17] W. Zhang. Reproducing SEAL. Blog post: <https://wtzhang99.github.io/blog/reproducing-seal/>, 2025.
- [18] Z. Zhang, G.P. Meyer, Z. Lu, A. Shrivastava, and A. Ravichandran. Knowledge distillation from VLM for long-tail visual recognition. *arXiv preprint arXiv:2408.16930*, 2024.
- [19] A. Zweiger, J. Pari, H. Guo, E. Akyürek, Y. Kim, and P. Agrawal. SEAL: Self-adapting language models. *arXiv preprint arXiv:2506.10943*, 2025.